

# Machine Learning in a hurry

Marc Lelarge  
marc.lelarge@inria.fr

February 6, 2024

## 1 Predictions with a statistical model

In machine learning, we have a **population** of  $N$  instances  $(x_i, y_i)_{i=1}^N$  but we first consider the statistical setting (called decision theory) where we know the **distribution** of  $(X, Y)$ . The question is then: given a sample  $x$  from this distribution, what is the best prediction for the associated  $y$ ?

### 1.1 Optimal prediction

**Modeling knowledge** as a probability distribution with a statistical model  $p(x, y)$ .

Prior for  $Y$ :  $p(y)$  in our case:  $p_0 = \mathbb{P}(Y = 0)$  and  $p_1 = 1 - p_0$ .

Generative model:  $p(x|y)$ . So that  $p(x, y) = p(x|y)p(y)$ .

We restrict ourselves to binary prediction, i.e. the target  $Y \in \{0, 1\}$ .

We denote by  $\hat{y}(X)$  the predictor for input  $X$ . For classification, we ask for  $\hat{y} : \mathbb{R} \rightarrow \{0, 1\}$ .

The **loss function** generalizes the natural notion of error:  $\text{loss}(\hat{y}, y) \in \mathbb{R}$ .

Ex: prediction error  $\text{loss}(\hat{y}, y) = \mathbb{1}(\hat{y} \neq y)$ .

#### Definition 1

The **risk** of a predictor  $\hat{y}$  is the expected loss:

$$R(\hat{y}) = \mathbb{E}[\text{loss}(\hat{y}(X), Y)] = \sum_{x,y} p(x, y) \text{loss}(\hat{y}(x), y).$$

#### Lemma 1

The predictor minimizing the risk is given by:

$$\hat{y}(x) = \mathbb{1} \left( p(1|x) \geq \frac{\text{loss}(1, 0) - \text{loss}(0, 0)}{\text{loss}(0, 1) - \text{loss}(1, 1)} p(0|x) \right).$$

**Proof.** Since

$$R(\hat{y}) = \sum_x p(x) \mathbb{E}[\text{loss}(\hat{y}(x), Y) | X = x],$$

we just need to compare the two terms:

$$\begin{aligned} \mathbb{E}[\text{loss}(0, Y) | X = x] &= \text{loss}(0, 0)p(0|x) + \text{loss}(0, 1)p(1|x) \\ \mathbb{E}[\text{loss}(1, Y) | X = x] &= \text{loss}(1, 0)p(0|x) + \text{loss}(1, 1)p(1|x), \end{aligned}$$

□

**Remark 1.** In the case where  $\text{loss}(0,0) = \text{loss}(1,1) = 0$  and  $\text{loss}(0,1) = \text{loss}(1,0) = 1$ , the optimal predictor is given by  $\hat{y}(x) = \arg \max_{y \in \{0,1\}} p(y|x)$ , which is the maximum a posteriori (MAP) rule.

Since our generative model is typically described with  $p(x|y)$ , we can rewrite the optimal predictor with Bayes rule:

$$\hat{y}(x) = \mathbb{1} \left( \frac{p(x|1)}{p(x|0)} \geq \frac{p_0 (\text{loss}(1,0) - \text{loss}(0,0))}{p_1 (\text{loss}(0,1) - \text{loss}(1,1))} \right).$$

## Definition 2

The likelihood ratio is defined as:  $\mathcal{L}(x) = \frac{p(x|1)}{p(x|0)}$  and a likelihood ratio test is a test of the form:  $\hat{y}(x) = \mathbb{1} (\mathcal{L}(x) \geq \eta)$  for some  $\eta > 0$ .

**Remark 2.** In the case where  $\text{loss}(0,0) = \text{loss}(1,1) = 0$  and  $\text{loss}(0,1) = \text{loss}(1,0) = 1$ , and  $p_0 = p_1$ , the MAP rule reduces to:  $\hat{y}(x) = \mathbb{1} (\mathcal{L}(x) \geq 1) = \arg \max_{y \in \{0,1\}} p(x|y)$ , which is the maximum likelihood (ML) rule.

## 1.2 Confusion matrix and ROC curve

The confusion matrix is given by:

	$Y = 0$	$Y = 1$
$\hat{y} = 0$	True negative	False negative
$\hat{y} = 1$	False positive	True positive

A few definitions:

- true positive rate (TPR):  $\mathbb{P}(\hat{y} = 1|Y = 1)$  also called sensitivity or recall.
- false negative rate (FNR):  $\mathbb{P}(\hat{y} = 0|Y = 1) = 1 - \text{TPR}$  also known as type II error.
- false positive rate (FPR):  $\mathbb{P}(\hat{y} = 1|Y = 0)$  also known as type I error.
- true negative rate (TNR):  $\mathbb{P}(\hat{y} = 0|Y = 0) = 1 - \text{FPR}$ .
- precision:  $\mathbb{P}(Y = 1|\hat{y} = 1) = \frac{p_1 \text{TPR}}{p_1 \text{TPR} + p_0 \text{FPR}}$ .
- $F_1$  score is the harmonic mean of precision and recall.

Note that, we have

$$\begin{aligned}
R(\hat{y}) &= p_0 (\mathbb{P}(\hat{y} = 1|Y = 0)\text{loss}(1,0) + \mathbb{P}(\hat{y} = 0|Y = 0)\text{loss}(0,0)) \\
&\quad + p_1 (\mathbb{P}(\hat{y} = 0|Y = 1)\text{loss}(0,1) + \mathbb{P}(\hat{y} = 1|Y = 1)\text{loss}(1,1)) \\
&= p_0 ((\text{FPR})\text{loss}(1,0) + (1 - \text{FPR})\text{loss}(0,0)) \\
&\quad + p_1 ((1 - \text{TPR})\text{loss}(0,1) + (\text{TPR})\text{loss}(1,1)) \\
&= \underbrace{p_0 (\text{loss}(1,0) - \text{loss}(0,0))}_{\alpha} \text{FPR} - \underbrace{p_1 (\text{loss}(0,1) - \text{loss}(1,1))}_{\beta} \text{TPR} \\
&\quad + \underbrace{p_0 \text{loss}(0,0) + p_1 \text{loss}(1,1)}_{\gamma}.
\end{aligned}$$

Since  $\alpha, \beta \geq 0$  and  $\gamma$  is a constant, there is a trade-off between  $\text{TPR} \uparrow$  and  $\text{FPR} \downarrow$ . This trade-off is captured by the receiver operating characteristic (ROC) curve corresponding to  $\max \text{TPR}$  as a function of  $\text{FPR}$  and can be captured by varying the loss function.


Indeed, note that we have  $R(\hat{y}) = \alpha \text{FPR} - \beta \text{TPR} + \gamma$  so that maximizing the TPR at a given FPR is the same as minimizing the risk  $R(\hat{y})$ . But we have shown that the optimal predictor minimizing  $R(\hat{y})$  is given by:  $\hat{y}(x) = \mathbb{1}(\mathcal{L}(x) \geq \frac{\alpha}{\beta})$ , hence the maximum TPR at a given FPR is still given by a likelihood ratio test. In other words, the ROC curve is obtained by varying the threshold  $\eta$  in the likelihood ratio test.

Setting  $\eta = 0$  or  $\eta = \infty$  corresponds to the two extreme points of the ROC curve:  $(0,0)$  and  $(1,1)$ . Also, for a given  $\pi \in [0,1]$ , the random predictor  $\hat{y}(x) = 1$  with probability  $\pi$  and 0 with probability  $1 - \pi$  corresponds to the point  $(\text{FPR}, \text{TPR}) = (\pi, \pi)$ . Hence the ROC curve is always above the diagonal. Finally, given two points on the ROC curve  $(\text{FPR}(\eta_1), \text{TPR}(\eta_1))$  and  $(\text{FPR}(\eta_2), \text{TPR}(\eta_2))$ , the point  $(t\text{FPR}(\eta_1) + (1-t)\text{FPR}(\eta_2), t\text{TPR}(\eta_1) + (1-t)\text{TPR}(\eta_2))$  is obtained by the random (suboptimal) predictor equal to  $\mathbb{1}(\mathcal{L}(x) \geq \eta_1)$  with probability  $t$  and to  $\mathbb{1}(\mathcal{L}(x) \geq \eta_2)$  with probability  $1 - t$ . Hence the ROC curve is concave.


### Proposition 1

*The points  $(0,0)$  and  $(1,1)$  are always on the ROC curve. The ROC curve is always above the diagonal and is concave.*

### 1.3 Example



**Figure 1:** Statistical model: Gaussian mixture with class 0 in blue and class 1 in orange



**Figure 2:** ROC curve associated to the Gaussian Mixture above